

CLOUD BASED FILE STORAGE USING NEXTCLOUD

EMANUELE RUOPPOLO –
SM3800049



GOALS OF THE PROJECT

- Implementing a file storage that ensures:

- **Broad network accessibility;**
- **Users' facilities;**
- **Security measures;**



- Using Nextcloud:
 - **Intuitive interface;**
 - **Full customizable in its features;**
 - **Highly compatible for the deployment in a Docker container for its setup;**

SETUP OF DOCKER CONTAINER

- Nextcloud has been deployed in a Docker container using a Docker-compose file:
 - A file `.yml` has been written with the needed specifics:
 - 2 volumes, one for Nextcloud one for a database;
 - 1 network for enabling communications between the volumes;
 - 2 services: Nextcloud, MariaDb, each with its specifics, for Nextcloud:
 - Port `'8080:80'` enabled between the container and the local host
 - Admin user and password set to login
 - Container created by running: ``docker compose up -d``
 - Container shut down by running ``docker compose down``
 - Nextcloud accessible from local browser typing ``http://localhost:8080`` and using admin credentials set in the compose file



NEXTCLOUD DEPLOYED FEATURES

Nextcloud allows for a full personal configuration, admins can easily install apps to guarantee many services in terms of security, accessibility and users' commodities. In our case these are the main features deployed:

- **Registration and authentication:** Nextcloud has a clear interface where it's easy login in, sign up and log out;
- **Authorizations management:** different roles within the system have been imposed, admin and users, each with different permissions. Users have also been grouped together to ease communication with them;
- **File management:** Users have an assigned Quota of storage, based on their needs, and they can manage this storage by uploading, downloading and deleting files, also a feature has been implemented that notices a user when its storage is almost full.

SECURITY MEASURES

To ensure users' security and isolation Nextcloud issues an access token for each client to use for any HTTP request and **encrypts** the clients' passwords in its database, it has a **Logging** page on its interface to check all the suspicious activities and allows for antiviruses that control the integrity of users' files.

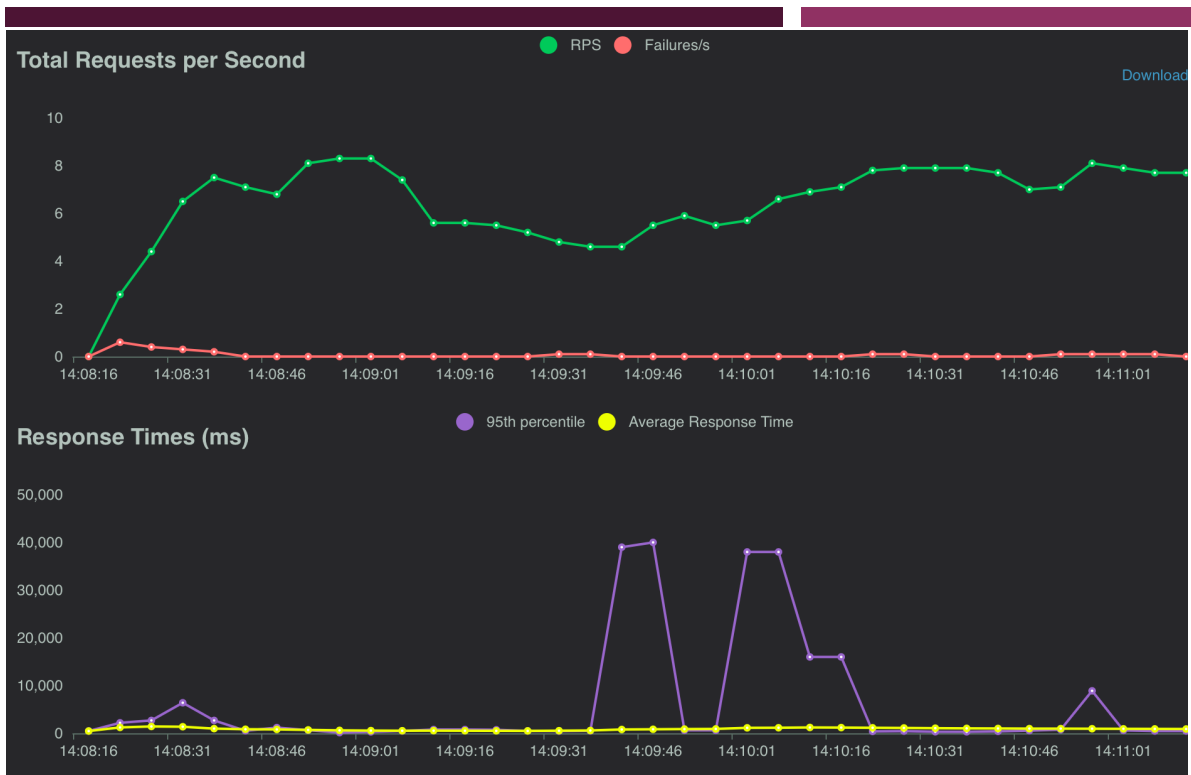
- **Registration** via the web interface has been made possible via e-mail-based sign-up with verification link;
- **Data encryption server-side** can be enabled;
- **Two factors authentication (2FA)** has been set in order to ensure the login;
- **Password recovery** via e-mail has been enabled;
- **Password constraints** such as minimum length, capital letters, special characters and numbers presence have been imposed

PERFORMANCES EVALUATION WITH LOCUST

- Locust stands as an open-source tool rooted in Python, dedicated to performance testing. It simulates vast user numbers to measure the scalability and performance of web applications;
- It allows to simulate scenarios where multiple user interactions are done simultaneously;
- A load is generated to rigorously test the application or service in question;

Procedure:

- Multiple users (90) have been created (and similarly deleted) with a docker dedicated command, to each of them have been assigned the same storage Quota (3GB);
- Files of different sizes (1KB, 1MB, 1GB, .jpg) have been created and added to a load data local repository;
- A `tasks.py` file has been created to make Locust available to test the system with different HTTP requests: **head, propfind, get, put, delete**.




30 users



90 users

RESULTS...

... AND FAILURES



180 users

EXPLORING SPAWN RATE AND FILE WEIGHT DEPENDENCIES

Two other tasks files have been made in order to evaluate the performances of the system when stressed in a certain scenario, like allowing more users logging in simultaneously uploading small size files, or less users uploading bigger files. The results are here presented:

Method	50%ile (ms)	66%ile (ms)	75%ile (ms)	80%ile (ms)	90%ile (ms)	95%ile (ms)	99%ile (ms)	100%ile (ms)
HEAD	18	22	25	30	44	3400	6200	7600
PROPFIND	21	26	30	35	55	1800	5800	7100
DELETE	31	33	34	35	40	330	6000	7700
PUT	36	42	47	51	73	870	4900	6800
GET	25	30	35	38	58	670	5500	8000
Aggregated	30	33	37	40	58	670	5800	8000

Table 1: Test performed for small files with 60 users, spawn rate = 10, for $\simeq 2$ min.

Method	50%ile (ms)	66%ile (ms)	75%ile (ms)	80%ile (ms)	90%ile (ms)	95%ile (ms)	99%ile (ms)	100%ile (ms)
HEAD	28	32	45	36	50	520	1300	1300
PROPFIND	33	46	51	53	56	59	1100	1200
DELETE	33	35	35	36	42	94	1500	2400
PUT	64	77	79	81	84	550	1600	1900
GET	39	52	60	61	63	65	1200	1200
Aggregated	37	49	56	61	79	120	1400	2400

Table 2: Test performed for big files with 20 users, spawn rate = 5, for $\simeq 2$ min.

SCALABILITY OF THE SYSTEM: ON-CLUSTER

The results have been obtained testing the system on a `small` machine. An idea for scaling the system in a much powerful environment could be deploy it **on-cluster**, a network that interconnects nodes with a network.

pros	cons
Full horizontal scalability	Complexity of the system set-up and effort in its maintenance
Data resilience due to built-in redundancy in case of failures	Costs of the infrastructure and its maintenance
Independent security policies	
Full control of the infrastructure	

SCALABILITY OF THE SYSTEM: ON-CLOUD

Another chance is to deploy the system on a cloud-based environment. There are many, one of these that seems reasonable is AWS S3, that provides a scalable infrastructure, ensuring data security and pay-as-you-go costs. In general these are the features of a on-cloud deployment:

pros	cons
On-demand scalability without hardware expenses	Data privacy due to the unknown management of the infrastructure
Global accessibility	Dependency to a third party, service outages
Pay-as-you go approach, minimizing costs and optimizing resource utilization	

CONCLUSIONS

- Nextcloud offers a user-friendly interface and useful services including robust **user authentication, role-based access control, private storage allocation, and comprehensive admin management capabilities**
- Nextcloud allows for the standards security measures of a cloud-based-storage;
- Locust tests demonstrate good performances in case of a small-environment deployment;
- Scalability can be addressed in either an on-cluster environment or a cloud-based one.